



DeepLearning.AI

Source Systems, Data Ingestion, and Pipelines

Week 3



DeepLearning.AI

DataOps

Week 3 Overview

DataOps


DataOps

Set of practices and cultural habits centered around building robust data systems and delivering high quality data products.

DevOps

Set of practices and cultural habits that allow software engineers to efficiently deliver and maintain high quality software products.

The Three Pillars of DataOps



Automation

**Continuous Integration
and Continuous
Delivery (CI/CD)**



Infrastructure as code

Code that acts to deploy the resources required to run your data pipeline



AWS CloudFormation



HashiCorp

Terraform

Week 3 Labs



Hands-on experience creating resources using Terraform



Monitoring Data Quality & other Observability Metrics

Interviews with Industry Experts



Chris Bergh




Barr Moses



Abe Gong



Chad Sanderson





DeepLearning.AI

DataOps

Chris Bergh
CEO at DataKitchen




DeepLearning.AI

DataOps - Automation

DataOps Automation


DevOps Automation

Continuous Integration and Continuous Delivery (CI/CD)



DataOps Automation


CI/CD can be applied directly to code and data within your data pipeline



DataOps Automation

Automation of data pipeline running


No automation: run all processes manually



DataOps Automation

Automation of data pipeline running

Pure scheduling: run stages of your pipeline according to a schedule




DataOps Automation

Automation of data pipeline running




A directed acyclic graph (**DAG**)



DataOps Automation

Key Underpinning CI/CD : Version control

Track changes in the **code**




**Revert to a
previous version**

DataOps Automation

Key Underpinning CI/CD : Version control

Track changes in the **code** and the **data**

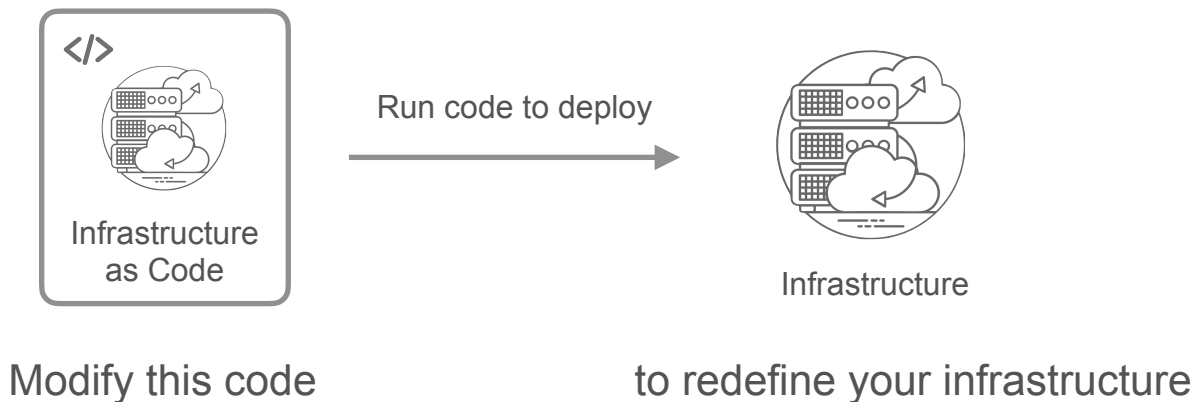


**Revert to
previous version**

DataOps Automation

Key Automation Aspect: Infrastructure as Code


Maintain the design of your infrastructure as a codebase



DataOps Automation

Version control over your entire infrastructure

Track changes in the **code** defining your infrastructure



**Revert to a
previous version**


DataOps

The Undercurrents




DataOps

The Undercurrents



DataOps

The Undercurrents





DeepLearning.AI

DataOps - Automation

Infrastructure as Code

Infrastructure as Code

1970s

**Configuration
Management**

Challenge : configuration of a series of physical machines



To automate some configuration tasks

Infrastructure as Code


1970s

Configuration
Management


2006



Amazon Elastic Compute
Cloud (Amazon EC2)



Infrastructure as Code



HashiCorp

Terraform




AWS CloudFormation




ANSIBLE

Infrastructure as Code



- Manage infrastructure resources on the cloud using code
- No manual clicking or writing bash scripts

Course 1 Week 2 Lab



AWS CloudFormation



HashiCorp

Terraform

Infrastructure as Code Tools



- Used in the lab
- Well supported with great documentation



AWS CloudFormation

- Native to AWS
- Well supported with great documentation

Terraform Configuration Files (C1W2 Lab)

```
1 resource "aws_s3_bucket" "data_lake" {
2   bucket_prefix = "${var.project}-datalake-${data.aws_caller_identity.current.account_id}-"
3 }
4
5 resource "aws_s3_bucket_public_access_block" "data_lake" {
6   bucket = aws_s3_bucket.data_lake.id
7
8   block_public_acls       = true
9   block_public_policy     = true
10  ignore_public_acls      = true
11  restrict_public_buckets = true
12 }
13
14 resource "aws_s3_bucket" "scripts" {
15   bucket_prefix = "${var.project}-scripts-${data.aws_caller_identity.current.account_id}-"
16 }
17
18 resource "aws_s3_bucket_public_access_block" "scripts" {
19   bucket = aws_s3_bucket.scripts.id
20
21   block_public_acls       = true
22   block_public_policy     = true
23   ignore_public_acls      = true
24   restrict_public_buckets = true
25 }
26
27 resource "aws_s3_object" "glue_job_script" {
28   bucket = aws_s3_bucket.scripts.id
```



```
1 resource "aws_glue_catalog_database" "analytics_database" {
2   name       = "${var.project}-analytics-db"
3   description = "Database for performing analytics on OLTP data"
4 }
5
6 resource "aws_glue_connection" "rds_connection" {
7   name = "${var.project}-rds-connection"
8
9   connection_properties = {
10     JDBC_CONNECTION_URL = "jdbc:mysql://${var.host}:${var.port}/${var.database}"
11     USERNAME            = var.username
12     PASSWORD            = var.password
13   }
14
15   physical_connection_requirements {
16     availability_zone = data.aws_subnet.private_a.availability_zone
17     security_group_id_list = [data.aws_security_group.db_sg.id]
18     subnet_id           = data.aws_subnet.private_a.id
19   }
20 }
21
22 resource "aws_glue_crawler" "s3_crawler" {
23   name = "${var.project}-analytics-db-crawler"
24   database_name = aws_glue_catalog_database.analytics_database.name
25   role         = aws_iam_role.glue_role.arn
26
27   s3_target {
28     path = "s3://${aws_s3_bucket.data_lake.bucket}/gold"
29   }
30 }
```



Terraform Configuration File (s3.tf)

Keyword Resource Type Name of the resource

```
1  resource "aws_s3_bucket" "data_lake" {
2      bucket_prefix = "${var.project}-datalake-${data.aws_caller_identity.current.account_id}-"
3  }
```

1. Set up the S3 bucket

```
4
5  resource "aws_s3_bucket_public_access_block" "data_lake" {
6      bucket = aws_s3_bucket.data_lake.id
7
8      block_public_acls      = true
9      block_public_policy    = true
10     ignore_public_acls     = true
11     restrict_public_buckets = true
12 }
```

Key-value pairs

2. Configure the bucket and provide public access to it

Terraform Language

```
1  resource "aws_s3_bucket" "data_lake" {
2    bucket_prefix = "${var.project}-datalake-${data.aws_caller_identity.current.account_id}-"
3  }
4
5  resource "aws_s3_bucket_public_access_block" "data_lake" {
6    bucket = aws_s3_bucket.data_lake.id
7
8    block_public_acls      = true
9    block_public_policy    = true
10   ignore_public_acls     = true
11   restrict_public_buckets = true
12 }
13
14 resource "aws_s3_bucket" "scripts" {
15   bucket_prefix = "${var.project}-scripts-${data.aws_caller_identity.current.account_id}-"
16 }
17
18 resource "aws_s3_bucket_public_access_block" "scripts" {
19   bucket = aws_s3_bucket.scripts.id
20
21   block_public_acls      = true
22   block_public_policy    = true
23   ignore_public_acls     = true
24   restrict_public_buckets = true
25 }
26
27 resource "aws_s3_object" "glue_job_script" {
28   bucket = aws_s3_bucket.scripts.id
```

Domain-Specific Language

HCL (HashiCorp Configuration Language)



HashiCorp

Terraform

The HCL Syntax

```
resource "aws_vpc" "main" {  
  cidr_block      = "10.0.0.0/16"  
  instance_tenancy = "default"  
  
  tags = {  
    Name = "main"  
  }  
}
```

Create or update a VPC



Virtual private
cloud (VPC)

```
resource "aws_instance" "web" {  
  ami          = data.aws_ami.ubuntu.id  
  instance_type = "t3.micro"  
  
  tags = {  
    Name = "HelloWorld"  
  }  
}
```

Create and Update an EC2 instance



Instance

The HCL Syntax

```
resource "google_compute_instance" "default" {  
  name          = "my-instance"  
  machine_type  = "n2-standard-2"  
  zone          = "us-central1-a"  
  
  tags = ["foo", "bar"]  
  
  boot_disk {  
    initialize_params {  
      image = "debian-cloud/debian-11"  
      labels = {  
        my_label = "value"  
      }  
    }  
  }  
}
```

Provision a GCP (Google Cloud Platform)
compute instance

Terraform Language

HCL is a declarative language

You just have to declare what you want the infrastructure to look like

```
1  resource "aws_s3_bucket" "data_lake" {
2    bucket_prefix = "${var.project}-datalake-${data.aws_caller_identity.current.account_id}-"
3  }
4
5  resource "aws_s3_bucket_public_access_block" "data_lake" {
6    bucket = aws_s3_bucket.data_lake.id
7
8    block_public_acls      = true
9    block_public_policy    = true
10   ignore_public_acls     = true
11   restrict_public_buckets = true
12 }
```



S3 Bucket


s3.tf

The desired end-state of the infrastructure:
resources and their configurations


Terraform is Highly idempotent!

Terraform Language

HCL is a declarative language




Create 5
EC2 Instances




Terraform Language

HCL is a declarative language




Create 5
EC2 instances



Imperative/Procedural Language



Imperative: specify the exact sequence of configuration tasks.





DeepLearning.AI

DataOps - Automation

Terraform - Creating an EC2 instance

Terraform

You write the configuration files to define your resources



Terraform prepares your workspace

1. Installs necessary files
2. Creates an execution plan

Terraform

You write the configuration files to define your resources



Terraform prepares your workspace




You approve the execution plan



Terraform applies the proposed steps

Terraform



Terraform

Use any IDE of your choice



Install Terraform in your environment



Use AWS credentials to authenticate Terraform

Developer / Terraform / Tutorials / AWS / Install


Install Terraform

9min | Terraform ▶ Video Interactive

Show Terminal



Reference this often? [Create an account](#) to bookmark tutorials.



To use Terraform you will need to install it. HashiCorp distributes Terraform as a [binary package](#). You can also install Terraform using popular package managers.



DeepLearning.AI

DataOps - Automation

Terraform - Defining variables and outputs



DeepLearning.AI

DataOps - Automation


**Terraform - Defining data
sources and modules**

Data Sources


Data Sources

Data blocks to reference resources created outside Terraform or in another Terraform workspace.

Data Sources



Data Sources





DeepLearning.AI


Lab Walkthrough

Implementing DataOps Automation with Terraform


Implementing DataOps with Terraform



- Architectural diagram of the resources
- Overview of the lab's steps




Architectural Diagram



Architectural Diagram

- Provided to you
- Defined as data blocks in Terraform configuration files




Architectural Diagram




SSH Key Pair:


- Public key
- Private key




Encrypt and decrypt
exchanged messages



Architectural Diagram




Architectural Diagram



Implementing DataOps with Terraform

Run terraform to create your resources






DeepLearning.AI

DataOps - Observability & Monitoring

Data Observability


Observability and Monitoring



DevOps Observability



Advent of the Cloud




Move toward Distributed Systems



Observability tools

Gain visibility into systems' health

Metrics



CPU and RAM usage



Response time

- Quickly detect anomalies
- Identify problems
- Prevent downtime
- Ensure reliable software products

Data Observability

**Monitor the health of
data systems**



Observability tools

**Monitor the health and
quality of data**



Health/quality of your data

What is High Quality Data?

High Quality Data

Accurate

Complete

Discoverable

Available in a timely manner

Low Quality Data

Inaccurate

Incomplete


Hard to find

Late

Represents exactly what stakeholders expect:

- well-defined schema
- data definitions

Quality of Data



Value



WORSE THAN

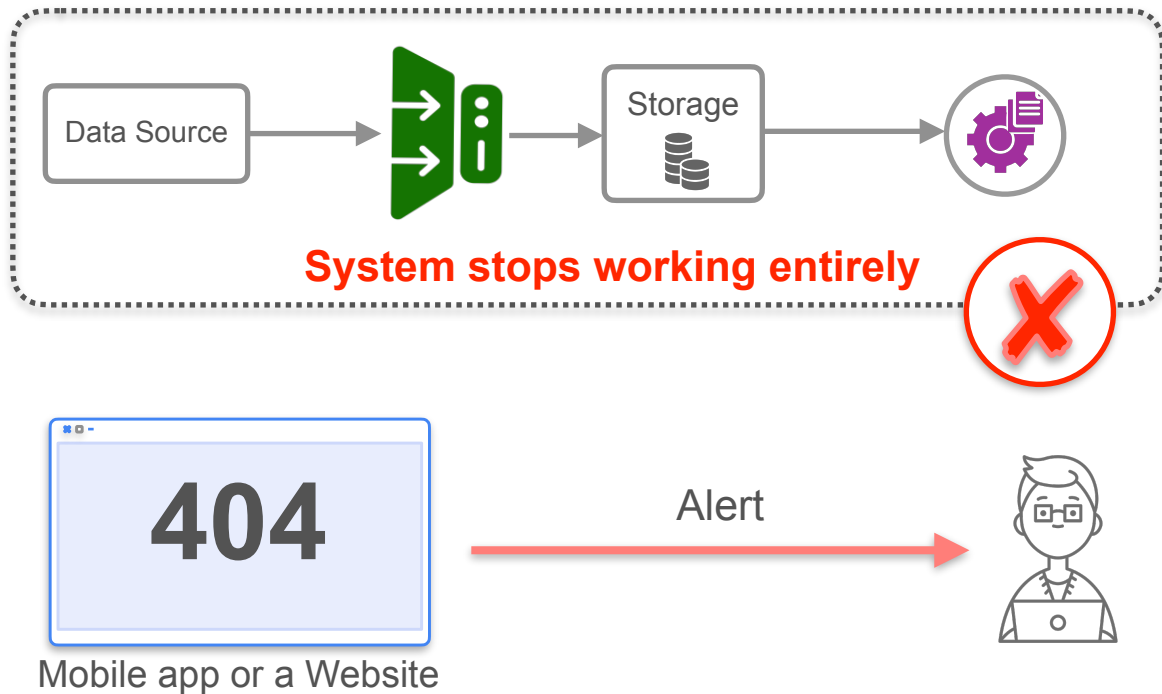


NO DATA



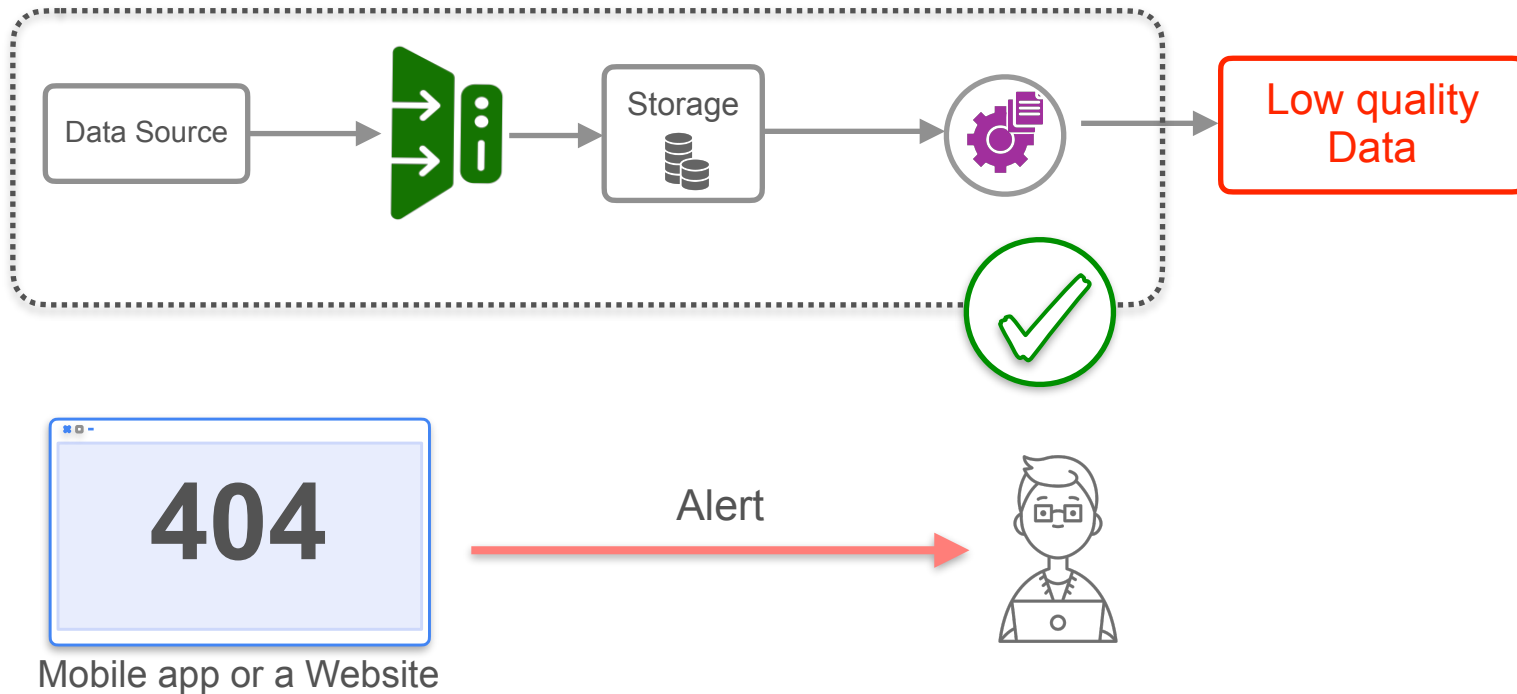
Data Observability

System failure is your best case scenario



Data Observability

Data suffers a breaking change




Scenario



U.S. based
company



Sell products in Europe



sale id	sale price (USD)
1	10
2	20
3	50

Total: 80



sale id	sale price (EUR)
1	9.15
2	18.31
3	45.77

Total: 73.23

It looks like there was
a sudden Drop in
Revenue



DeepLearning.AI

DataOps - Observability & Monitoring

Barr Moses

Co-founder & CEO of Monte Carlo




DeepLearning.AI

DataOps - Observability & Monitoring

Monitoring Data Quality

Data Quality Metrics

- **Volume:** Total number of records ingested in each batch or over some time interval



- **Null values:** Total number of null values in a table

col 1	col 2	col 3
1	10	100
2	null	null
null	50	600

- **Distribution:** Range of values in a particular column stays within some thresholds

col 1	col 2	col 3
1	10	100
2	20	300
3	50	600

- **Freshness:** The *difference* in time between *now* and the *timestamp of the most recent* record in your data



- Identify and focus on the most important metrics
- Avoid creating confusion and “alert fatigue”

Data Quality Metrics



Stakeholders

What do stakeholders care about most for this use case?

Stakeholder Requirement

Current data: no more than 24 hours old

What to monitor?

“Freshness” of the data: measure when the latest records were ingested

Data Quality Metrics

Stakeholders interested in product sales revenue

Sales

sale id	product id	purchase amount
1	2	10
2	3	20
3	4	null

Product

product id	product name	SKU
1	product 1	yk3
2	product 2	yk3
3	product 3	sj6


What to monitor?

- Verify that all sales records were successfully ingested
- No null values in sales record

What may be less important?

- Product SKU numbers match product descriptions
- Each customer's postal code was recorded correctly

Testing Ingested Data



- Build in checks or tests to verify that the schema and data types stay consistent.
- Identify problems as early as possible before they propagate further down your pipelines.



DeepLearning.AI

DataOps - Observability & Monitoring

Abe Gong


Co-founder & CEO of Great Expectations




DeepLearning.AI

DataOps - Observability & Monitoring

Great Expectations - Core Components




1. Specify the data
2. Define your expectations
3. Validate your data against the expectations




Great Expectations Workflow

- **Data Context:** entry point for the Great Expectations API
- **Great Expectations API:** classes and methods that allow you to connect data sources, create expectations and validate your data




Great Expectations Workflow

- **Data source:** SQL Database, a local file system, an S3 bucket, a Pandas DataFrame
- **Data asset:** collection of records within a data source
 - Table in a SQL database
 - File in a file system
 - Join query asset
 - Collections of files matching a pattern



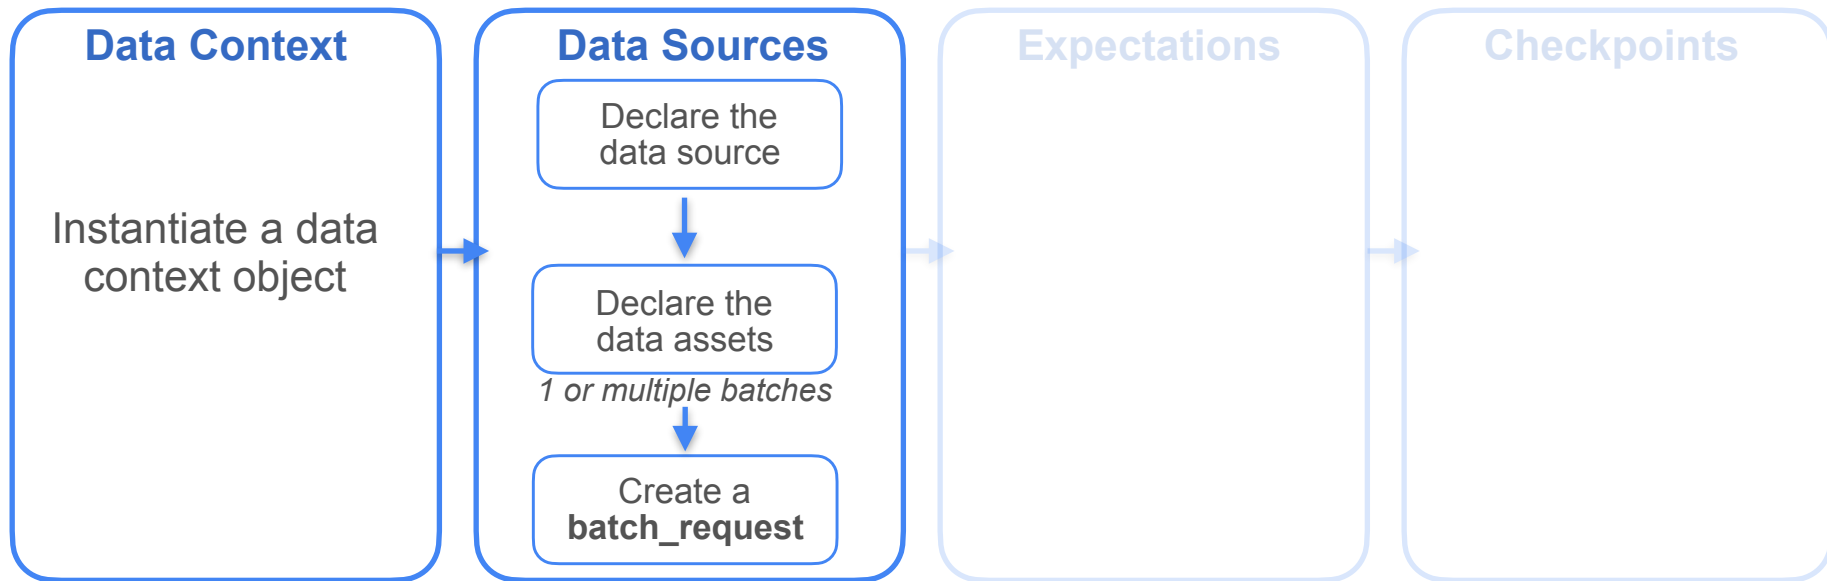
Great Expectations Workflow

- **Data source:** SQL Database, a local file system, an S3 bucket, a Pandas DataFrame
- **Data asset:** collection of records within a data source
- **Batches:** partitions from your data asset
 - partitions by date
 - partitions by column values




Great Expectations Workflow

- **Data source:** SQL Database, a local file system, an S3 bucket, a Pandas DataFrame
- **Data asset:** collection of records within a data source
- **Batches:** partitions from your data asset
- **Batch_request:** primary way to retrieve data from the data asset




Great Expectations Workflow

- **Expectation:** statement that you can use to verify if your data meets a certain condition
 - `expect_column_min_to_be_between`
 - `expect_column_values_to_be_unique`
 - `expect_column_values_to_be_null`



Great Expectations Workflow


- **Validator:** expects a `batch_request` and its corresponding expectation suite
 - Manual interaction
 - Or automating the validation process using a checkpoint object



Great Expectations Workflow

Metadata is saved in backend stores:

- Expectation Store
- Validation Store
- Checkpoint Store
- Data docs





DeepLearning.AI

DataOps - Observability & Monitoring

Great Expectations - Workflow Example


Workflow Example

DVD rental Database


Local postgresSQL database

payment
payment_id
customer_id
staff_id
rental_id
amount
payment_date

- payment_id contains unique IDs
- customer_id does not contain null values
- all the values in the amount column are non-negative



Workflow Example





DeepLearning.AI

DataOps - Observability & Monitoring

Chad Sanderson
CEO of Gable.ai



DeepLearning.AI

DataOps - Observability & Monitoring

Amazon CloudWatch

Amazon CloudWatch



System Level Metrics:

- CPU utilization
- Disk I/O
- Network traffic
- Memory usage

Amazon CloudWatch



Made aware when
metrics reflect issues



CloudWatch Alarms

System Level Metrics:

Provide a general understanding of how your resources are performing

- CPU utilization
- Disk I/O
- Network traffic
- Memory usage

Custom Metrics:

Allow you to monitor specific aspects of your application that are not covered by default system metrics

- Number of transactions processed
- Response time of an API endpoint
- Number of active users

- Define thresholds for these metrics
- Establish a baseline
Measure performance of your system under different loads and conditions
- CloudWatch retains metrics data for up to 15 months

Common Metrics for RDS

CPU Utilization

- High value: your RDS instance might be under heavy load
- Values over 80-90%: can lead to performance bottlenecks

RAM Consumption

- High RAM consumption: can slow down performance

Disk Space

- Value consistently above 85%: you may need to delete or archive data to free up some space

Database Connections

- The number of active connections to your database
- Number of connections approaching the maximum limit: can lead to connection errors and application failures




DeepLearning.AI

DataOps

Summary

DataOps Summary



DataOps Summary



Chris Bergh



Barr Moses



Abe Gong



Chad Sanderson

In This Week's Labs



HashiCorp

Terraform

- A cloud agnostic tool for developing and maintaining your data infrastructure as a codebase
- Infrastructure as Code tools help make deployments consistent and repeatable



Amazon CloudWatch



great
expectations